

Applied Machine Learning

Course number: W207


Prof. Alexander I. Iliev, Ph.D.

Applied Machine Learning

Lecture 13 ...


- *Recommendation Systems*
- *PCA, SVD*
- *CNN*

Signal Processing



Signal Processing

- Convolving two signals:



box vs box

spiky vs box


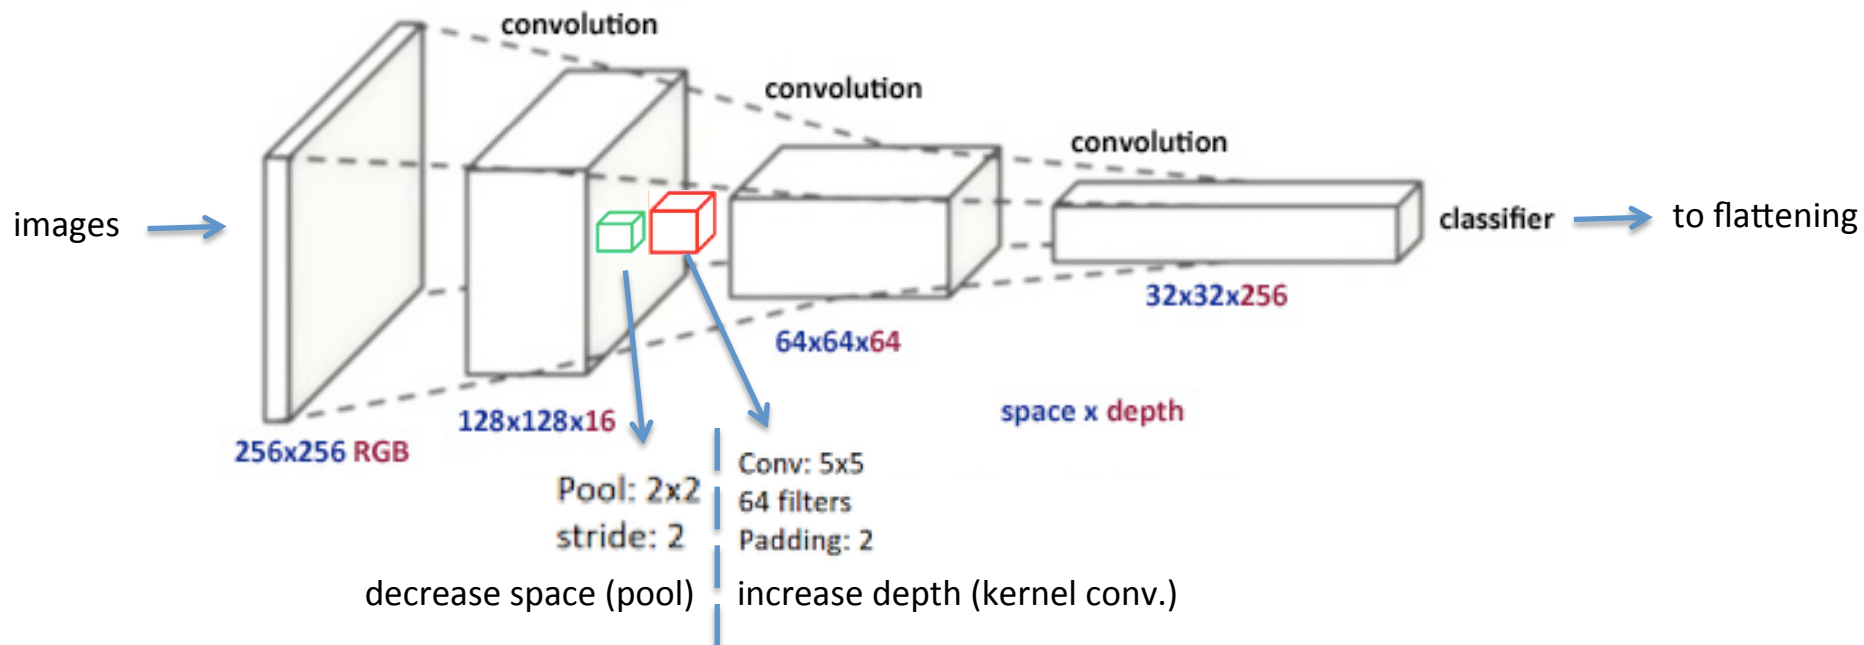


Image recognition - CNN


- In NN **matrixes are multiplied (dot product)** by other matrixes all the way through the network
- We start with **bigger images (space)** and after every layer they become more, but **smaller (depth)**



Model

Labels + Images
mapping

Pre-processing of our data



Batch

Hidden Layer 1
(Convolution 1)

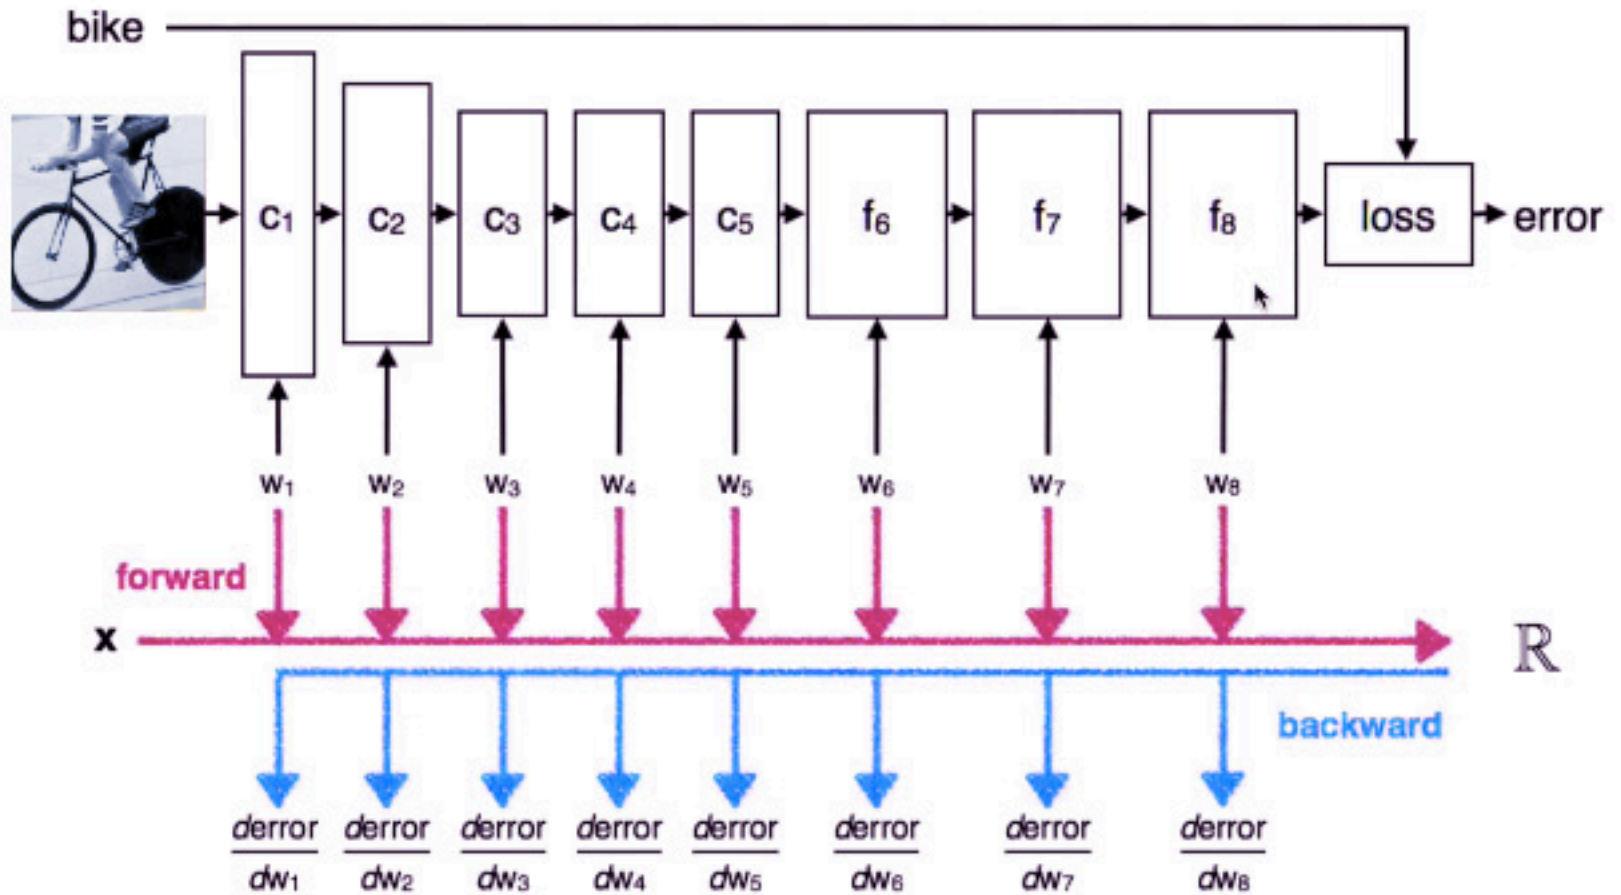
Hidden Layer 2
(Convolution 2)

Hidden Layer 3
(Fully Connected)

Training

Forward propagation

Model



Backward propagation



CNN clarifications

- These terms mean the same thing in CNN:

weights = weight matrix = feature matrix = feature map = kernel = filter

- We can specify filters by taking specific parts of images for example or we can use pre-specified filters
- We then convolve these filters with other images and search for suitable overlap, where parts that are not interesting will = 0 and those that are interesting will > 0. Higher values are better !!!
- For that we use ReLU and is called activation. It is applied to every single pixel of the image.

CNN clarifications

- The **neurons** in each layer **perform the same mathematical operation**
- We **feed** it with **the data from the previous layer** in the network.
- **Dropout** provides random neuron dropout **to prevent overfitting**
- Fully connected layer is using **softmax** for probability (last layer)
- **Softmax will give us a probability** of our findings [0:1] based on the classes that are more likely to occur.
- We pick the maximum probability class using **np.argmax()**

CNN clarifications

- We **train** the network **using Gradient Descent = Backpropagation**
- The **number of hidden layers and neurons** is determined **experimentally**
- **Flip a convolutional network** and you **get a De-convolutional** network where you can create an image out of text
- **Use 2D and 3D CNN** any time we have **spatial data** (where the space and order matter), such as: **images, sounds, text**, etc.
- Remember: for **regression** we use the **MSE** or **RMSE** and for **classification** we use the **softmax** function

CNN clarifications

- How to improve the performance of a CNN?
 - Vary all hyperparameters such as:
 - Number of inputs (decrease large number of pixels to all images)
 - Number of layers
 - Type, frequency and % of pooling
 - Dropout to prevent overfitting
 - Image pre-processing techniques:
 - Use uniform size and aspect ratio for all images
 - Use scaling, cropping and padding of images to make them of equal size
 - Use mean pixel values across all training examples (for each pixel)
 - Image normalization: subtract the mean from each pixel and then dividing the result by the STD
 - Convert images to grayscale (only if it makes sense)
 - Determine stride step for the kernel (filter) at each step
 - To increase performance, consider holding on pooling for several steps:
`conv2d > activation > conv2d > activation > max_pool`
This way you use more intermediate features initially to maximize the effect of training at the cost of computational efficiency